Poisoned Identifiers Survive LLM Deobfuscation: A Case Study on Claude Opus 4.6

Luis Guzmán Lorenzo

Model tested: Claude Opus 4.6 (1M context)

2026-03-22 through 2026-04-05

Abstract

When an LLM deobfuscates JavaScript, can poisoned identifier names in the string table survive into the model's reconstructed code, even when the model demonstrably understands the correct semantics? Using Claude Opus 4.6, we found that explicit verification prompts ("verify each name matches the math") failed to prevent propagation (12/12), while task reframing ("write from scratch") substantially reduced propagation on the physics artifact (from 100% to 20%) and eliminated it on the pathfinding artifact (0%). In our tested setting, the implicit frame of the task changed naming accuracy where explicit instructions did not.

We tested this across 192 repeated inference runs on two code archetypes (a force-directed graph simulation and an A* pathfinding algorithm, 50 conditions, N=3–6) and observed three consistent patterns within the tested conditions:

1. **Wrong decoded identifiers persisted.** Poisoned names appeared in the model's deobfuscated code in every baseline run tested on both artifacts (physics: 8/8; pathfinding: 5/5). Matched controls showed this extends to terms with zero semantic fit (`combustion` for repulsion, `invoice` for heuristic) when the string table does not form a coherent alternative domain.
2. **Persistence coexisted with correct semantic commentary.** In 15/17 runs on the physics artifact, the model wrote wrong variable names in code while correctly describing the actual operation in comments (manually scored, unblinded, single author; see Section 3.3 for limitations).
3. **Task framing changed persistence.** Reframing from "deobfuscate this" to "write a fresh implementation" reduced propagation from 100% to 0–20% on the physics artifact (N=5) and to 0% on the pathfinding artifact (N=5), while an algorithmic consistency check confirmed the generated implementations preserved the checked structure (Appendix F). Explicit verification instructions had no effect (12/12 across 4 prompt variants).

Matched-control experiments showed that terms with zero semantic fit persist at the same rate when the string table does not form a recognizable alternative domain. Per-term variation observed in earlier domain-gradient experiments is confounded with domain-level coherence and recoverability (Section 6).

These observations are from two code archetypes on one model family (Opus 4.6 primary; Haiku 4.5 limited spot-check). Broader model and archetype generalization is needed.

1 Introduction

When Claude Opus 4.6 deobfuscates JavaScript containing poisoned identifier names, explicit verification instructions ("verify each name matches the math") do not prevent the wrong names from

appearing in the output (12/12 runs). Reframing the task as generation ("write a fresh implementation from scratch") substantially reduces propagation while preserving the checked algorithmic structure (Appendix F). The implicit frame of the task changed naming accuracy where explicit instructions did not.

Here is what the output looks like:

```
attraction: 4000,           // Repulsion force strength
----------                     ---------------------
      |                                |
 identifier: WRONG              comment: CORRECT

amplification: 0.92,        // Velocity damping (friction)
-------------                  ----------------------
      |                                |
 identifier: WRONG              comment: CORRECT
```

Figure 1: The dual-representation pattern (Phase B, R1 run 1, Opus 4.6). Wrong identifier in code, correct description in comment. Observed in 15/17 Phase B runs where the primary endpoint was positive (at least one poisoned identifier in code blocks). Scored manually by the first author without blinding; see Section 3.3.

There are two ways to read this. Source recovery: decoded string-table names are the closest available ground truth, and a well-calibrated model should use them. Alternatively: the model does not apply a capability it shows in other settings (name correction on readable code, Section 5.1), even when explicitly told to.

A pure inability-to-verify account does not explain these results well, though the data do not isolate the underlying mechanism. The strongest defensible claim is: task framing changes naming behavior without altering the checked algorithmic structure.

Scope. This is a case study on two code archetypes (a JavaScript force-directed graph simulation and an A* pathfinding algorithm) tested primarily on Claude Opus 4.6 (Haiku 4.5 limited spot-check). Phase A (exploratory, N=1) used Claude Code CLI with tool use. Phase B (replicated, N=3–6) used the Anthropic Messages API without tool use. Results may reflect model-specific or interface-specific behaviors. We do not claim generality beyond these tested conditions.

2 Related Work

2.1 Code obfuscation and deobfuscation

Pre-LLM baselines. Lu, Coogan, and Debray [17] describe semantics-based automatic deobfuscation of JavaScript using dynamic analysis and slicing (2011). Their approach recovers functionally equivalent code without requiring knowledge of the obfuscation scheme, establishing the pre-LLM baseline for automated JS deobfuscation. The transition from such deterministic approaches to LLM-based deobfuscation introduces the identifier-naming question central to this paper: deterministic tools recover whatever literals the code contains, while LLMs reconstruct names that may reflect training priors, decoded string tables, or both.

CASCADE [1]: Google's hybrid LLM + compiler deobfuscator. Uses Gemini for prelude function identification (99.56% accuracy via few-shot prompting), then hands off to JSIR for deterministic constant propagation, with sandboxed V8/QuickJS for dynamic execution. Achieves 98.93% success on standard obfuscator.io output. CASCADE's JSIR constant propagation recovers

the original string literals (before obfuscation encoded them), not semantically meaningful names derived from algorithmic analysis. If the original names were poisoned before obfuscation (our attack model), we predict but did not test that JSIR would recover those poisoned names faithfully, since it restores original literals regardless of semantic correctness.

webcrack [4]: Purpose-built AST-level deobfuscator for javascript-obfuscator output. Supports string array decoding (base64, RC4), control flow unflattening, and dead code elimination via pattern matching. Issue history indicates failure cases under some transform combinations (e.g., CFF + dead code injection, issue #44). Our RC4-encoded production build was not fully decoded by webcrack's automated pipeline in Phase A testing.

JS-Confuser [6]: Alternative obfuscator using state-machine CFF (explicit state transitions instead of dispatch strings). We expect webcrack cannot reverse this pattern (untested; webcrack's source matches only the split-string dispatch pattern).

2.2 Identifier naming and LLM code understanding

Wang et al. [10] measure the impact of misleading variable names on LLM code analysis tasks (using CodeBERT-family models), finding that shuffled or misleading names degrade performance more than random strings because they actively misdirect the model. This supports the general premise that identifier names influence LLM code processing, though the setting (analysis tasks on code models) differs from our deobfuscation setting.

"The Code Barrier" [11] introduces semantics-preserving obfuscations including cross-domain term substitution, finding that removing meaningful names severely degrades intent-level tasks. Their cross-domain renaming experiments provide relevant background for our domain-gradient observations. Our contribution is specific to the deobfuscation string-table setting: we observe that wrong-but-plausible names are actively preserved in reconstructed code, not merely that they degrade understanding.

CoTDeceptor [5] demonstrates that multi-stage obfuscation strategies can evade CoT-enhanced LLM vulnerability detectors, in line with the general finding that adversarial content can bypass model analysis.

2.3 Obfuscation vs. LLM code analysis

Li et al. [12] categorize obfuscation techniques (layout, data flow, control flow) against LLM vulnerability detection, finding both degradation and, in some cases, unexpected improvements. **"Digital Camouflage"** [13] tests Claude, Gemini, and GPT-4o against compiler-level obfuscation (CFF, bogus control flow) on LLVM-obfuscated C functions, confirming cross-model degradation. **FORGEJS** [14] benchmarks LLM vulnerability detection in JavaScript and finds generally unreliable performance. **JsDeObsBench** [16] provides the first systematic benchmark for LLM JavaScript deobfuscation, evaluating GPT-4o, Mixtral, Llama, and DeepSeek-Coder across obfuscation types from variable renaming to control flow flattening. Its focus is on overall deobfuscation quality (syntax accuracy, execution reliability, code similarity); our work addresses a different question: whether specific poisoned identifiers survive into the model's output. **Yefet et al.** [7] establish adversarial examples for code models. These studies show that obfuscation frequently degrades LLM code analysis under tested settings. Our contribution extends this in a specific direction: surface-level poisoning (wrong names in the string table) not only degrades understanding but can produce contradictory outputs where the model shows correct semantic knowledge and incorrect identifier usage in the same response.

2.4 Tool-augmented LLM agents

SWE-bench [15] and related work on agentic coding workflows show that tool use (file read/write, shell execution) produces qualitatively different behavior than single-shot inference. Our observation that progressive analytical escalation appeared in Claude Code (Phase A, with tools) but not in API testing (Phase B, without tools) is in line with this (Section 7.1).

Historical precedent. 1988 paragraph-book anti-piracy [2, 3]: see Section 8.3.

3 Method

3.1 Stimuli

Exploratory vs. hypothesis-testing. Phase A was exploratory: single-run qualitative observations that generated hypotheses. Phase B sub-phases 0–3 were hypothesis-testing with replication. Sub-phases 4–6 were targeted follow-ups, designed to test specific predictions (translation-frame, domain-boundary stability) and boost sample sizes on key conditions from earlier sub-phases. We did not preregister hypotheses. The same author designed stimuli, ran experiments, and scored results. We know these are limitations. The automated scoring pipeline and published raw outputs help, but do not fully compensate.

Two phases of experimentation:

Phase A (exploratory, N=1): 17 pill designs (we use "pill" as shorthand for individual adversarial test stimuli, after the poison-pill metaphor) tested in Claude Code CLI sessions (with tool use). 7 production build variants. 28 tests. Results are qualitative observations that generated hypotheses.

Phase B (systematic, N=3–6): 50 conditions tested via Anthropic Messages API (no tool use). 192 automated runs across 2 models (183 Opus, 9 Haiku) in 10 sub-phases (0–9). Results are quantitative with per-condition replication. All conditions at N=3 or higher. Sample sizes preclude formal inferential statistics; we report observed proportions and treat consistency across varied conditions as the primary evidential basis.

Control: pill 09 (clean code, standard obfuscation, no adversarial content).

Unit of analysis. The 192 Phase B runs test varied conditions on two code archetypes (a force-directed graph simulation, primary; an A* pathfinding algorithm, replication). The design supports claims about identifier propagation behavior under varied prompt, obfuscation, and naming conditions on these archetypes. It does not support claims about LLM deobfuscation in general. We frame the paper as a case study accordingly.

3.2 Protocol

Phase A: Fresh Claude Code session from `/tmp`. File deleted before each session. Model's complete response recorded manually. Follow-up prompts used in some sessions to test bypass resistance. Tool use (Bash, file read/write) available to the model.

Phase B: Automated via Anthropic Messages API. Each pill's code sent as a user message with the exact prompt text. No tool use. Full response, token counts, and stop reason recorded as structured JSON. Automated scoring with code-block-level analysis (see 3.3).

Temperature. Most runs used API default temperature. A targeted comparison at temperature 0 (N=5, Phase 5) showed identical 100% propagation for core terms (`attraction` 5/5, `amplification` 5/5), confirming the effect is deterministic for high-fit names and not a sampling artifact.

3.3 Metrics

Primary endpoint: Whether a poisoned identifier is preserved in code blocks of the model's output (binary, per term per run). This is the unit of analysis for all propagation claims.

Secondary endpoint (dual-representation): Whether wrong identifier in code AND correct semantic description in comments/prose appear in the same response (binary, per run). This pattern appeared in 15/17 Phase B runs where the primary endpoint was positive.

Scoring for this endpoint: a run was counted as dual-representation positive when (a) at least one poisoned identifier appeared in a code block, AND (b) the same response contained a comment or prose description that correctly named the actual operation (e.g., "Repulsion force strength" adjacent to `attraction`). "Correct description" was defined as any phrase that would allow a reader to identify the true physical operation without reference to the variable name. This was scored manually by the first author on all 17 applicable Phase B runs. Scoring was not blinded; the scorer knew which names were poisoned. Ambiguous descriptions (e.g., "force parameter" without specifying direction) were excluded from the positive count. The 15/17 rate should be read with this unblinded, single-scorer limitation in mind.

Tertiary endpoints:

- **Refusal:** Keyword detection for "malware," "I won't," "do not run," "compromised," etc. (binary per run). We manually validated a subset of 10 runs across conditions and confirmed the keyword proxy matched human judgment in all checked cases (Appendix D).
- **Semantic corruption scoring:** For each poisoned name, check whether it appears in code blocks (`"‘...”‘`) vs. only in prose. Scoring definitions (applied per poisoned term per run):
  - **Propagated:** Wrong name appears in code blocks, correct name does not appear in code blocks.
  - **Flagged:** Both wrong and correct names appear in code blocks.
  - **Corrected:** Correct name appears in code blocks, wrong name does not.
  - **Absent:** Neither wrong nor correct name appears in code blocks.

  N.B.: Naive full-text scoring produces false negatives because the model writes correct physics descriptions in comments alongside wrong variable names. This motivated the code-block-level scoring. A model-based adjudication check (Appendix E) validated this scorer against an independent Haiku 4.5 blind adjudicator: 17/20 agreement (85%), with all disagreements in the conservative direction (the automated scorer undercounted propagation).
- **Time/tokens:** API-reported processing time and token usage.
- **Functional reconstruction:** Binary, tested in Phase A production builds only.

3.4 Experimental Matrix

Phase A (Claude Code CLI): 28 additional exploratory runs, N=1 per condition.

3.5 Limitations

- Two models tested (Opus 4.6, Haiku 4.5). Broader generalization requires GPT-4o, Gemini, open-source models.

Sub-phase	Condition group	N (runs)	Model(s)
Phase 0	Pill 10 replication	5	Opus
Phase 0	Pill 03 replication	3	Opus
Phase 0	Pill 10 cross-model	3	Haiku
Phase 0	Pill 03 cross-model	3	Haiku
Phase 1	Prompt variants (4 prompts)	12	Opus
Phase 1	Domain gradient (4 domains)	8	Opus
Phase 1	Obfuscation gradient (3 levels)	6	Opus
Phase 1	Task frames (4 frames)	8	Opus
Phase 2	Layer gradient (5 layers)	10	Opus
Phase 2	Canary	3	Opus
Phase 3	Extended domains (gamedev, overlap)	6	Opus
Phase 3	Multi-agent verify	6	Both
Phase 3	Task frame replications	6	Opus
Phase 4	Framing experiment (3 frames)	15	Opus
Phase 4	N-boost: gradient, canary, multi-agent	20	Opus
Phase 5	N-boost: 4 domains + heavy obf + temp 0	21	Opus
Phase 6	N-boost: obf gradient, task frames, layer gradient	27	Opus
Phase 7	Matched lexical controls (2 terms, full table)	10	Opus
Phase 8	Full matched-table control (all terms, no fit)	5	Opus
Phase 9	Second artifact (A* pathfinding): baseline, framing, control	15	Opus
Total	**50 conditions**	**192 runs**	**183 Opus + 9 Haiku**

- Two code archetypes tested (force-directed graph, A* pathfinding). Further archetype diversity is needed for broader generalization.
- Phase A: single run per condition, manual interpretation.
- Phase B: 3–6 runs per condition, automated scoring.
- Self-designed stimuli with implicit bias toward known model behaviors.
- Phase A and Phase B use different interfaces (Claude Code with tools vs. raw API without tools). Behaviors observed in Phase A may be specific to tool-using workflows.
- The scoring methodology (code-block extraction + regex) may miss propagation in non-standard response formats.
- Most Phase B runs used API default temperature.

Negative results and non-replications are reported alongside positive findings throughout this paper.

4 Taxonomy of Adversarial Code-Semantic Manipulations

We tested five categories of adversarial content embedded in identical JavaScript code (a force-directed graph simulation). Four produced no effect on the model's analysis; one produced consistent propagation. The negative results establish specificity.

4.1 Ineffective Manipulations

Four categories produced no effect: inert textual content (wrong comments, legal threats, CBRN strings, combinations; pills 01–05), and attention capture via malware decoys (pill 14, Phase A only). The model identified all inert content as "dead code" or "prompt injection" in 7/7 tests. See Appendix A for the pill-by-pill index.

4.2 Structural Executable Content: Triggers Refusal

Functions with their own coherent malware-like logic (not physics code with malware names, but actual malware-shaped operations).

Tested: Functions building credential objects, ransom notes, exploit lists (pill 03). Realistic functions with actual API calls (pills 17, production builds v4–v7).

Observed: The model refused to assist with these functions on readable code. On obfuscated code, refusal reversed: the model entered "forensics mode" and fully deobfuscated the malware (pill 13). On partially obfuscated code with realistic malware, the model classified the entire file as malware and refused to read further.

Key condition: Refusal requires that the malware functions have their own internal logic that is not reclassifiable as benign. Functions named `credential_harvest` that compute $\sqrt{\Delta x^2 + \Delta y^2}$ are identified as "physics with edgy names." Functions that call `require('child_process').execSync(...)` cannot be reclassified.

Phase B refusal rates (API, replicated):

Condition	Model	N	Refusal rate
Pill 03 (malware functions, readable)	Opus 4.6	3	**100%**
Pill 03 (malware functions, readable)	Haiku 4.5	3	**100%**
Pill 10 (poisoned names, no malware)	Opus 4.6	5	0%
Pill 10, security framing	Opus 4.6	5	**100%**
Pill 10, practical framing	Opus 4.6	5	20%
Pill 14 (malware + poisoned)	Opus 4.6	5	**100%**
Pill 22-4layer / 22-5layer	Opus 4.6	5 each	**100%**

Refusal is binary and replicates across Opus and Haiku (6/6). Security-framed prompts trigger refusal even without malware.

4.3 Naming-Semantic Manipulation: Propagates After Obfuscation

Renaming physics variables to wrong-but-plausible-in-domain terms: `repulsion` to `attraction`, `damping` to `amplification`, etc.

Tested on readable code: Pills 06, 07, 08. The model caught inverted semantics by tracing mathematical signs (pill 06), caught swapped node/edge variables by reading object shapes (pill 07), and propagated truncated constants without correction (pill 08: TAU=6.283 used as-is with "~2pi" annotation).

Tested after obfuscation: Pill 10 (testbed, 130 lines). All wrong names propagated into the model's deobfuscated output with the dual-representation pattern described in Section 1 (Figure 1).

Key condition (revised by Phase 8 matched controls): Initial experiments suggested wrong names needed to be plausible in the target domain. Matched-control tests (Section 6.1) showed a broader pattern: decoded names persist even with zero semantic fit (`combustion` for repulsion) when the replacement table lacks a coherent alternative-domain signal. Malware-domain

names (`exfil_range`) are still corrected; the model re-derives from algorithm structure (pill 11), likely because malware terms trigger a distinct recognition pathway.

4.4 Numerical Precision Manipulation: Propagated in All Tested Conditions

Truncated mathematical constants: $\tau \approx 6.283$ (real $2\pi = 6.28318\ldots$), $1/\sqrt{2} \approx 0.707$ (real $= 0.70710\ldots$).

Observed: The model propagated truncated values without correction in all conditions where we explicitly tested it: readable code (pill 08, Phase A) and obfuscated code (pill 10, Phase A+B). It added an approximate annotation ("~2pi") but did not substitute the precise value. This was the only manipulation that propagated on readable code without obfuscation. Systematic replication of constant propagation was not prioritized in Phase B, where the focus was on identifier names.

5 Identifier Persistence Under Deobfuscation

5.1 Observation

On readable code, the model corrected wrong names using cues consistent with algorithm pattern recognition, object field shapes, mathematical sign tracing, domain-name cross-reference, and numerical precision. Wrong names were overridden in every readable-code test.

After obfuscation (control flow flattening, RC4 string encoding), the observed outputs showed decoded names dominating reconstructed identifiers. The model produced the dual-representation pattern (Figure 1): wrong identifiers in code, correct descriptions in comments.

5.2 Persistence Under Verification Instructions

In Phase B, we tested four prompt variants on pill 10 (N=3 per variant, 12 total runs). Prompts ranged from baseline ("deobfuscate this") to explicit adversarial warning ("the string table may contain adversarial names; cross-check each decoded name against the algorithm structure").

Result: 12/12 runs propagated the wrong names in code regardless of prompt:

Prompt	`attraction` in code	`amplification` in code
A: "deobfuscate this" (baseline)	**3/3**	**3/3**
B: "deobfuscate + verify names"	**3/3**	**3/3**
C: "verify each name matches the math"	**3/3**	**3/3**
D: "string table may be adversarial, cross-check"	**3/3**	**3/3**

The model spent more time on verification prompts (52.6s vs 36.6s baseline) but still propagated. This is difficult to explain as a simple failure to check.

5.3 The Framing Experiment

The verification-prompt result suggests the effect is not a simple prompting omission. But what kind of effect is it? One explanation that fits the data: deobfuscation may activate a translation-like workflow where decoded string-table entries serve as "source text." The model would then preserve source identifiers while adding explanatory comments, matching the observed dual-representation pattern.

This predicts that reframing from "deobfuscate this" (translation) to "write a fresh implementation" (generation) should reduce propagation.

We tested this prediction (Phase 4, N=5 per frame):

Frame	`attraction`	`repulsion`	`amplification`	`damping`
"Deobfuscate" (translation)	**5/5**	1/5	**5/5**	5/5
"Rewrite with correct names"	4/5	5/5	3/5	5/5
"Write fresh from scratch"	**1/5**	**5/5**	**0/5**	**5/5**

The generation frame reduced `attraction` propagation from 100% to 20% and `amplification` from 100% to 0%. The model produced correct names (`repulsion`, `damping`) at 100% when freed from the translation frame. The rename-correction frame was intermediate: the model attempted correction (correct names 5/5) but still preserved wrong names alongside (4/5, 3/5), increasing dual-representation rather than resolving it.

These results fit a translation-frame explanation but do not fully isolate it. The generation frame also changes fidelity constraints, output freedom, and the cost of renaming. An algorithmic consistency check (Appendix F) confirmed that generation-frame outputs preserve the checked structure on the physics artifact (5/5 runs) and on pathfinding (5/5 runs). This rules out drift on the checked components as an explanation for correct naming. What the data do establish: the identifier-persistence effect is not an intrinsic inability. The model uses correct names readily under a different task frame.

5.4 Propagation Across Conditions

Baseline propagation (Phase B, code-block-level scoring):

Condition	Model	N	`attraction` in code	`amplification` in code
Pill 10 (physics-domain poison)	Opus 4.6	5	**5/5**	**5/5**
Pill 10 (physics-domain poison)	Haiku 4.5	3	**3/3**	2/3
Pill 24 (canary comment)	Opus 4.6	3	2/3	0/3

A research-context canary comment partially reduces propagation for marginal-plausibility terms; core terms still propagated 2/3.

Obfuscation gradient (pill 10 source, 3 levels):

Level	`attraction` in code	`amplification` in code	N
Light (base64, CFF 0.5)	**5/5**	**5/5**	5
Medium (base64, CFF 0.5, self-defending)	**5/5**	**5/5**	5
Heavy (RC4, CFF 0.75, split strings)	**6/6**	0/6	6

Light obfuscation is sufficient to trigger propagation (16/16 for `attraction` across all levels). `amplification` dropped to 0/6 at heavy obfuscation, where split-string encoding fragments the name across chunks. This recoverability interaction is one reason the domain-gradient experiments (Section 6.2) should be treated as exploratory.

6 What Determines Whether Wrong Names Are Corrected?

6.1 Matched Controls (Phase 7–8)

The domain-gradient experiments (Section 6.2) suggested per-term semantic fit determines propagation, but this was confounded with domain-level coherence. To isolate the variable, we created length-matched terms with zero semantic fit: `combustion` (10 chars, like `attraction`) for repulsion, `communication` (13 chars, like `amplification`) for damping. In one condition only 2 terms were changed; in another, ALL config property names were replaced with length-matched no-fit terms.

Condition	All terms replaced?	Domain coherent?	Propagation
pill-10 (physics terms)	Yes	Yes (physics)	FLAGGED 5/5
pill-25 (2 no-fit terms, rest physics)	No	Mostly	FLAGGED 5/5
pill-27 (ALL no-fit terms, incoherent)	Yes	No	**FLAGGED 5/5**
pill-20-engineering (ALL engineering)	Yes	**Yes (engineering)**	**CORRECTED 5/5**

FLAGGED = both wrong and correct names appeared in code blocks; the wrong name was preserved. CORRECTED = only correct names in code.

`combustion` persisted in code at the same rate as `attraction`, even when every term in the string table lacked semantic fit. Full correction occurred in the engineering coherent-domain condition (0/5 propagation), though other coherent domains (e.g., finance) produced mixed rather than clean correction.

The same pattern appeared on the second artifact (A* pathfinding): `invoice` (zero fit to heuristic estimation) persisted at 5/5, while the generation frame corrected all terms to proper pathfinding vocabulary (5/5). See Section 6.3.

6.2 Domain-Gradient Experiments (Phases 1–5, exploratory)

Earlier domain-pill experiments, which replaced all terms with coherent domain-specific alternatives, showed per-term variation:

Term	For operation	Domain	Propagation	N
`attraction`	repulsion	physics	100%	8
`decay`	damping	overlap	100%	3
`velocity`	repulsion	game dev	75%	8
`yield`	repulsion	finance	60%	5
`acceleration`	damping	game dev	0%	8
`inductance`	repulsion	engineering	0%	5

This looked like per-term semantic evaluation at first. However, the matched-control results (Section 6.1) showed this interpretation is confounded: the domain pills replaced ALL terms with coherent alternatives, creating a recognizable domain signal. The per-term variation likely reflects the model's confidence in the alternative domain, not per-term semantic evaluation. The `acceleration` result (0/8 despite being a core physics term) may reflect domain-level recognition of the game-dev replacement set rather than recognition that acceleration contradicts what damping does.

The intermediate points (e.g., yield 3/5) are underpowered and should not be over-interpreted. This initial gradient motivated the matched-control experiments (Section 6.1) that clarified the picture.

6.3 Second Artifact Replication (Phase 9)

To test whether these patterns are specific to the physics simulation, we repeated the core protocol on an A* pathfinding algorithm (~110 lines, graph-search domain). The pathfinding artifact uses a different algorithmic structure (priority-queue search vs. force integration), different domain vocabulary, and a different poison map:

Poisoned name	Correct name	Operation
`penalty`	`heuristic`	distance estimation method
`adjacentCost`	`diagonalCost`	diagonal movement cost
`closedSet` / `openSet`	(swapped)	A* frontier sets
`fScore` / `gScore`	(swapped)	path cost components
`ancestors`	`neighbors`	adjacent-cell expansion (same function, wrong label)

Scoring used the same code-block extraction methodology as the physics artifact. The generation-frame prompt was adapted to reference pathfinding rather than physics. Three conditions at N=5 each (15 runs total):

Condition	Physics artifact	Pathfinding artifact
Baseline (poisoned names)	FLAGGED 5/5	FLAGGED 5/5
Generation frame	CORRECTED 4–5/5	CORRECTED 5/5
Matched no-fit control	FLAGGED 5/5	FLAGGED 5/5

The pathfinding no-fit control used `invoice` (7 chars, like `penalty`) for the heuristic method and `purchaseCost` (12 chars, like `adjacentCost`) for diagonal movement cost.

Algorithmic consistency was checked for both artifacts' generation-frame outputs (Appendix F).

All three patterns replicated on the second artifact: persistence, framing-dependent correction, and no-fit-term persistence. This suggests the observed patterns are not unique to the first artifact, though two archetypes remain a limited basis for generalization.

7 Secondary Findings

7.1 Adversarial Salience Escalation (Phase A only)

In Phase A (Claude Code with tools), defense sophistication was inversely correlated with defense effectiveness. Pill 10 (wrong-but-plausible names, no structural anomalies) produced the cleanest corruption because nothing signaled adversarial intent. Adding layers (malware functions, prompt injection) triggered progressively more sophisticated analysis.

Phase B did not replicate this pattern. In the API (without tools), refusal is a step function at the malware-function boundary, not a gradient. The model either analyzes routinely or refuses. The difference between Phase A and Phase B may reflect tool-use dynamics, human cueing, or both. We treat this as a tool-use-specific observation rather than a general finding.

7.2 Cost Inflation (Phase A, Claude Code with tools)

As an engineering illustration: the undefended control (pill 09) was deobfuscated in 2.5 min with near-perfect output. The defended production build (v7) required 13.5 min for structural recovery and 3+ hours for a full reconstruction attempt (137 tool invocations, 268K tokens) that produced

non-functional output with poisoned parameter names throughout. These are single-run Phase A observations and should not be read as controlled effect-size estimates.

7.3 Multi-Agent Verification Pilot (Phase 3)

A fresh model instance was given the deobfuscated code (readable, with wrong names and correct comments) and asked to verify each name.

Model	N	Wrong names flagged	Fully corrected in code
Opus 4.6	3	2/3	0/3
Haiku 4.5	3	2/3	0/3

Both models detected the discrepancy but neither consistently removed the wrong names from code output. At N=3 per model, this is a pilot observation. The readable-code setting is cognitively different from deobfuscation, so the two effects may not share a common mechanism.

8 Discussion

8.1 Conjecture

One hypothesis that fits the data: during deobfuscation, decoded identifier names (high apparent reliability) may dominate over algorithmic patterns (degraded by obfuscation) in the model's output.

The matched-control results (Section 6.1) suggest a more specific hypothesis: the model may not evaluate decoded names individually but instead assess whether the full set of decoded names forms a coherent domain signal. When it recognizes a coherent alternative domain (engineering terms on physics code), it shifts vocabulary; when the names are incoherent or mixed, it defaults to preserving the decoded entries as given. This points toward in-context domain recognition rather than per-token semantic evaluation.

The framing result (Section 5.3) adds a second dimension: the preservation appears tied to the translation-like workflow specifically, since the generation frame eliminates it while preserving algorithmic structure.

These are behavioral descriptions, not established mechanisms. The data do not isolate the underlying cause.

8.2 Practical Effect of the Defense

The defense is a cost multiplier that degrades automated semantic recovery. It does not prevent a determined attacker from understanding the code's architecture. After 13.5 minutes, the model in our Phase A production-artifact test correctly identified: force-directed graph, 12 project nodes, 50 concept nodes, Canvas 2D rendering, spring/repulsion physics, particle system, camera system, and animation loop.

What the defense disrupted (in our observations): zero-effort casual analysis, correct parameter naming, and quick functional reproduction.

Practical implication for LLM-assisted workflows: Our framing result suggests that a generation-frame pass ("reimplement this from scratch") after deobfuscation may correct poisoned identifiers where verification prompts do not. This is testable in production toolchains.

8.3 Cost Multiplication, Not Protection

A browser-delivered artifact is ultimately inspectable. The defense adds cost to the LLM-assisted path but does not close the manual path. The framing parallel is to 1988 paragraph-book anti-piracy techniques [2, 3], which worked because the reader had no independent verification channel.

9 Future Work

1. **Tool-use vs. single-shot inference.** Adversarial salience escalation appeared in Claude Code (agentic, tool-using) but not via the API. Does tool use change how models allocate verification effort during deobfuscation? A controlled comparison (same model, same code, tool-use enabled vs. disabled) would isolate this.

2. **Model comparison.** GPT-4o, Gemini 2.5 Pro, and open-source models remain untested. The propagation endpoints (Section 6) are the highest-priority cross-model test.

3. **Custom obfuscator.** State-machine CFF (anti-webcrack) and polymorphic generation. We implemented a prototype environment-aware decoder in v7 but did not systematically test it in Phase B.

4. **Temporal stability.** Re-run key conditions after model updates.

5. **Constants propagation on readable code.** TAU=6.283 propagated without obfuscation in Phase A (pill 08), making it the only manipulation effective on readable code. This was not systematically replicated in Phase B and may be a more practical attack vector than identifier poisoning.

6. **Mapping the domain-recognition boundary.** Phase 8 matched controls showed that per-term semantic fit is confounded with domain-level coherence. Testing what makes a set of replacement terms "recognizably from another domain" (number of coherent terms needed, domain distance, term familiarity) would sharpen this finding.

7. **Multi-agent mitigation.** Initial testing (Section 7.3, N=3 per model) showed partial detection but 0/6 full correction. Testing alternative configurations (e.g., giving the verifier only pseudocode, forcing name generation from scratch) is the natural next step.

8. **Cross-model framing experiment.** The framing manipulation (Section 5.3) was tested only on Opus 4.6. Replicating it on Haiku, GPT-4o, and Gemini would test whether the translation-frame effect is model-specific or general. (The algorithmic consistency check for generation-frame outputs is completed and reported in Appendix F.)

9. **Competing string tables.** Providing the model with two conflicting sets of decoded names (e.g., a second string table containing correct names alongside the poisoned one) should force explicit conflict resolution rather than silent propagation. This would test whether the preservation pattern reflects lexical-source dominance or a more general default.

10 Conclusion

The headline finding is the framing asymmetry: explicit verification instructions did not reduce identifier propagation (12/12), but reframing from "deobfuscate" to "write fresh" substantially

reduced it on physics (100% to 20%) and eliminated it on pathfinding (0%), while preserving the checked algorithmic structure (Appendix F). For LLM-assisted code analysis workflows, this suggests that how a task is framed may matter more than how carefully the model is instructed to verify its output.

The matched-control experiments turned up a second finding we did not expect: terms with zero semantic fit (`combustion` for repulsion, `invoice` for heuristic) persisted at the same rate as plausible terms when the string table did not form a coherent alternative domain. Correction occurred in the engineering coherent-domain condition but not uniformly across other coherent domains. The initial per-term semantic-fit gradient is confounded with domain-level coherence and recoverability.

What remains unresolved: the mechanism. The data are compatible with default preservation of decoded string-table entries, modulated by domain-level coherence recognition and task framing, but they do not isolate the cause. The dual-representation pattern (wrong names in code, correct descriptions in comments, 15/17 on the physics artifact) is suggestive but manually scored and unblinded. The critical next steps are cross-model replication (GPT-4o, Gemini), additional code archetypes, and independent rescoring of the dual-representation endpoint.

Conflict of Interest

The first author operates kieleth.com, whose production JavaScript artifact is used as a case study in Appendix G. The controlled experiments (Sections 4–8) use synthetic test stimuli independent of the production artifact.

Data Availability

Raw JSON outputs for all Phase B runs (192 total), test stimuli (pills), batch configurations, and the automated scoring pipeline are available at: https://github.com/Kieleth/obfuscated-sentinel (Kieleth, kieleth.com)

Phase A session transcripts are not included (manual sessions, not machine-recorded).

Responsible Disclosure

We disclosed these findings to Anthropic prior to publication. The behaviors documented here are observable through normal API usage and do not exploit any vulnerability in access controls or safety systems.

References

[1] S. Jiang, P. Kovuri, D. Tao, and Z. Tan. "CASCADE: LLM-Powered JavaScript Deobfuscator at Google." *arXiv preprint arXiv:2507.17691*, 2025.

[2] *Sentinel Worlds I: Future Magic* paragraph booklet. Electronic Arts, 1988. Archived at mocagh.org/ea/futuremagicaus-paragraphs.pdf.

[3] J. Dott. "Wasteland." *The Digital Antiquarian*, 2016. filfre.net/2016/02/wasteland/.

[4] j4k0xb. webcrack: AST-level JavaScript deobfuscator. github.com/j4k0xb/webcrack.

[5] H. Li, M. Li, J. Zuo, et al. "CoTDeceptor: Adversarial Code Obfuscation Against CoT-Enhanced LLM Code Agents." *arXiv preprint arXiv:2512.21250*, 2025.

[6] MichaelXF. JS-Confuser: JavaScript obfuscator. github.com/MichaelXF/js-confuser.

[7] N. Yefet, U. Alon, and E. Yahav. "Adversarial Examples for Models of Code." *Proc. ACM Program. Lang.* (OOPSLA), 4, 2020. arXiv:1910.07517.

[8] D. Zhou, L. Ying, H. Chai, and D. Wang. "From Obfuscated to Obvious: A Comprehensive JavaScript Deobfuscation Tool for Security Analysis." In *NDSS*, 2026. arXiv:2512.14070.

[9] OWASP. "LLM Top 10 2025: LLM01 — Prompt Injection." genai.owasp.org/llmrisk/llm01-prompt-injection/.

[10] Z. Wang, L. Zhang, C. Cao, N. Luo, X. Luo, and P. Liu. "How Does Naming Affect LLMs on Code Analysis Tasks?" 2024. arXiv:2307.12488.

[11] S. L. Nikiema, J. Samhi, A. K. Kaboré, J. Klein, and T. F. Bissyandé. "The Code Barrier: What LLMs Actually Understand?" *arXiv preprint arXiv:2504.10557*, 2025.

[12] X. Li, Y. Li, H. Wu, et al. "A Systematic Study of Code Obfuscation Against LLM-based Vulnerability Detection." *arXiv preprint arXiv:2512.16538*, 2025.

[13] E. Böke and S. Torka. "Digital Camouflage: The LLVM Challenge in LLM-Based Malware Detection." *arXiv preprint arXiv:2509.16671*, 2025.

[14] Q. Fei, X. Liu, S. Li, et al. "Large Language Models Cannot Reliably Detect Vulnerabilities in JavaScript." (FORGEJS). *arXiv preprint arXiv:2512.01255*, 2025.

[15] C. E. Jimenez, J. Yang, A. Wettig, et al. "SWE-bench: Can Language Models Resolve Real-World GitHub Issues?" In *ICLR*, 2024. arXiv:2310.06770.

[16] G. Chen, X. Jin, and Z. Lin. "JsDeObsBench: Measuring and Benchmarking LLMs for JavaScript Deobfuscation." In *CCS*, 2025. arXiv:2506.20170.

[17] G. Lu, K. Coogan, and S. Debray. "Automatic Simplification of Obfuscated JavaScript Code: A Semantics-Based Approach." Technical Report, University of Arizona, 2011.

A Experiment Index

The full experimental matrix is in Section 3.4. Phase B totals: 192 automated runs (183 Opus, 9 Haiku) across 50 conditions in 10 sub-phases. Phase A: 28 exploratory runs. **Grand total: 220 runs.**

Raw JSON outputs for all Phase B runs: `experiments/results/phase{0..9}/`.

B Environment

Phase A (Claude Code CLI)

- Model: Claude Opus 4.6 (1M context)
- CLI: Claude Code v2.1.86 through v2.1.91

ID	Technique	State	Refusal	Names propagated	Time
Pill 01	Wrong comments	Readable	No	N/A	Instant
Pill 02	Legal threats (strings)	Readable	No	N/A	Instant
Pill 03	Malware functions	Readable	Partial	N/A	Instant
Pill 04	CBRN system prompts	Readable	No	N/A	Instant
Pill 05	All combined	Readable	No	N/A	Instant
Pill 06	Inverted semantics	Readable	No	0/4	Instant
Pill 07	Cascading rename	Readable	No	0/5	Instant
Pill 08	Trojan constants	Readable	No	TAU propagated	Instant
Pill 09	Clean (control)	Obfuscated	No	0 (correct output)	2.5 min
Pill 10	Poisoned names	Obfuscated	No	All (~10)	1.75 min
Pill 11	Malware-domain names	Obfuscated	No	0 (re-derived)	57 sec
Pill 12	License text	Obfuscated	No	N/A	~2 min
Pill 13	Pill 03 obfuscated	Obfuscated	No (forensics)	N/A	~3 min
Pill 14	Hybrid (03+10)	Obfuscated	No	Physics under-analyzed	8 min
Pill 15	Dual implementation	Readable	No	N/A	Instant
Pill 16	Decision graph	Readable	No	N/A	Instant
Pill 17	Realistic malware	Partial obf	Yes (alarm)	Not tested	~1 min
Prod v4	IIFE + decoder poison + decoy	Full obf	Yes	Not reached	2 min
Prod v7	Nested layers + hidden poison	Full obf	Yes / broke on push	All 20 params	13.5 min
Prod v7 recon.	Same, "reconstruct fully"	Full obf	N/A	Poisoned throughout	3+ hours

- Platform: macOS Darwin 24.6.0
- Subscription: Claude Max
- Test directory: /private/tmp (fresh sessions, no project context)
- Obfuscator: javascript-obfuscator v4.x with RC4 encoding

Phase B (Anthropic Messages API)

- Models: Claude Opus 4.6 (`claude-opus-4-6`),
  Claude Haiku 4.5 (`claude-haiku-4-5-20251001`)
- SDK: @anthropic-ai/sdk v0.82.0
- Max tokens per response: 16,384
- Temperature: not explicitly set; API defaults were used
- No `top_p`, `top_k`, or seed parameters set (API defaults)
- No tool use (single-shot inference)
- Each run: fresh API request with no conversation state, identical model ID, identical parameters, no cached context
- Automated scoring via code-block extraction + regex matching
- Test harness: `experiments/testbed/run-experiment.js`

C Exemplar Output

See Figure 1. Full response transcript:
`experiments/results/phase0/R1/pill-10-obfuscated-poisoned_claude-opus-4-6_run1.json` Pattern appeared in 15/17 Phase B runs where the primary endpoint was positive (at least one poisoned identifier in code blocks).

D Refusal Metric Validation

We manually reviewed 10 runs sampled across 4 condition families (pill 03 refusal, pill 10 deobfuscation, T4-B security framing, S1 layer gradient with malware). The keyword-based refusal detector agreed with human judgment in all 10 cases (5 true positives, 5 true negatives). No borderline cases were observed in the sample.

E Model-Based Scoring Adjudication (Phase 7)

A model-based blind adjudicator (Haiku 4.5, a different model from the generator) evaluated 20 Phase B responses for presence of `attraction` and `amplification` in code blocks. Agreement with the automated regex scorer: 17/20 (85%). All 3 disagreements were in the same direction: the adjudicator detected `attraction` in code where the regex missed it. The automated pipeline is conservative (undercounts propagation). No cases where automation inflated counts.

Scope: This validates the automated propagation scorer (the primary endpoint). It does not independently validate the manual 15/17 dual-representation count (the secondary endpoint), which relies on a single unblinded author scorer. Independent human rescoring of the dual-representation endpoint remains an open validation gap.

Results: `experiments/results/blind-scoring/blind-scoring-results.json`
Script: `scripts/blind-score.js`

F Algorithmic Consistency Check

We verified that generation-frame outputs ("write fresh from scratch") preserve the scored algorithmic components and key parameter values on both artifacts.

Physics artifact (outputs from Phase 4, checked during Phase 7; N=5 per frame):

Frame	Params (4)	Formulas (4)	Correct names	Wrong names
Generation	5/5 all 4	5/5 all 4	4/5	1/5
Translation	5/5 all 4	5/5 all 4	0/5	5/5

Checks: inverse-square repulsion, Hooke's law springs, velocity damping operation, center gravity, parameter values (4000, 0.004, 0.92, 0.0003).

Pathfinding artifact (Phase 9, N=5 per frame):

Frame	Key values (3)	Algorithmic checks (4)	Correct names	Wrong names
Generation	5/5 all 3	5/5 all 4	5/5	0/5
Translation	5/5 all 3	5/5 all 4	0/5	5/5

Key values: diagonal cost (1.414), max iterations (1000), obstacle ratio (0.25). Algorithmic checks: Manhattan heuristic present, open/closed set structure, g-score/f-score computation, neighbor expansion.

On both artifacts, the scored algorithmic components and key parameter values were identical across frames. Name correction under the generation frame is not a side effect of drift on the scored components and key parameter values.

Script: `scripts/score-equivalence.js`
Results (physics): `experiments/results/phase4/_equivalence_check.json`
Results (pathfinding): `experiments/results/phase9/`

G Production Case Note (Phase A)

We applied the defense to the kieleth.com graph.js production artifact (1,500 lines, 690KB after obfuscation combining malware-logic functions, poisoned physics names, truncated constants, and standard obfuscation). In Phase A testing (Claude Code with tools, N=1 per prompt), the model's initial response shifted from semantic deobfuscation to malware classification. Full reconstruction (3+ hours, 137 tool invocations) produced non-functional output with poisoned parameter names throughout. This is a single-artifact observation that illustrates the controlled findings at production scale but does not constitute independent replication.